

Alex Rivera

AI Engineer

alex.rivera@example.com · San Francisco, CA · github.com/alexrivera · linkedin.com/in/alexrivera

SUMMARY

AI engineer specializing in production ML systems: LLM platforms, RAG, agents, fine-tuning, and MLOps. Focused on reliable inference, rigorous evaluation, and measurable product impact.

EXPERIENCE

Senior AI Engineer · Nova Labs

2023 — Present · San Francisco, CA

- Lead LLM platform (RAG, agents, evals) powering core features for 50k+ users.
- Designed multi-tenant hybrid retrieval; improved answer faithfulness by 22%.
- Mentored 3 engineers; established model review and canary deployment practices.

Machine Learning Engineer · Horizon AI

2021 — 2023 · Remote

- Shipped vision + NLP document understanding; reduced manual review by 60%.
- Built K8s training pipelines with experiment tracking and regression tests.
- Defined offline/online evaluation loops with product partners.

ML Engineer · DataPulse

2019 — 2021 · New York, NY

- Productionized ranking and recommendation models for high-throughput APIs.
- Owned feature pipelines and monitoring for drift and latency SLOs.

SELECTED PROJECTS

Aether — Multi-agent research system

Agents with tools, memory, and eval gates; ~70% faster literature review.

VectorForge — Production RAG

Hybrid search + re-rank; 2M+ queries/mo at p99 < 180ms retrieval.

SignalNet — Anomaly detection

Streaming ML for infra metrics; 40% fewer false-positive pages.

TuneKit — LoRA toolkit

Shared fine-tuning CLI/evals adopted by 5 product teams.

SKILLS

ML & LLMs: PyTorch, Transformers, LoRA/PEFT, RLHF/DPO, RAG, Agents, vLLM, eval harnesses

MLOps: Kubernetes, Docker, MLflow/W&B, CI/CD for models, GPU optimization

Software: Python, TypeScript, FastAPI, PostgreSQL, Redis, AWS/GCP

EDUCATION

Carnegie Mellon University · 2017 — 2019

M.S. Machine Learning

University of California, Berkeley · 2013 — 2017

B.S. Computer Science